



Recurrent Neural Networks

Richard Dirauf, M.Sc.

Machine Learning and Data Analytics (MaD) Lab

Friedrich-Alexander-Universität Erlangen-Nürnberg (FAU)

MLTS Exercise, 26.01.2024

- How can we utilize neural networks to process sequential data?
- Passing complete sequences to a conventional network is computationally very expensive
 - long training
 - gradients converge very slowly
 - Temporal context gets lost
- **Solution:** Adapt the time related characteristic into the architecture of the net


Input: Harry Potter invented a new spell.

Dictionary:

a	1
aaron	2
...	...
harry	2,039
...	...
potter	6,453
...	...
spell	8,940
...	...
zoo	10,000


Input: Harry Potter invented a new spell.

Dictionary:




Recurrent Neural Networks (RNN)

MLTS Exercise 10




$$h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t + b_h)$$
$$y_t = \sigma(W_{hy}h_t + b_y)$$

σ : sigmoid / softmax



The RNN Cell Up Close

MLTS Exercise 10



Hidden state: $h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t + b_h)$


W_{hh} : Weights of previous hidden state

W_{xh} : Weights of the input

b_h : Bias for state update

The RNN Cell Up Close

MLTS Exercise 10



Output:


$$y_t = \sigma(W_{hy}h_t + b_y)$$

W_{hy} : Weights of current hidden state

b_y : Bias for output update

Backpropagation Through Time (BPTT)

MLTS Exercise 10



- The backward pass takes activations off the stack to compute the error derivatives at each time step
- Backpropagate through time all the way to the initial states to get the gradient of the cost function with respect to each initial state

(BPTT) Calculate The Cost

MLTS Exercise 10

Total cost is the sum of losses at each time step

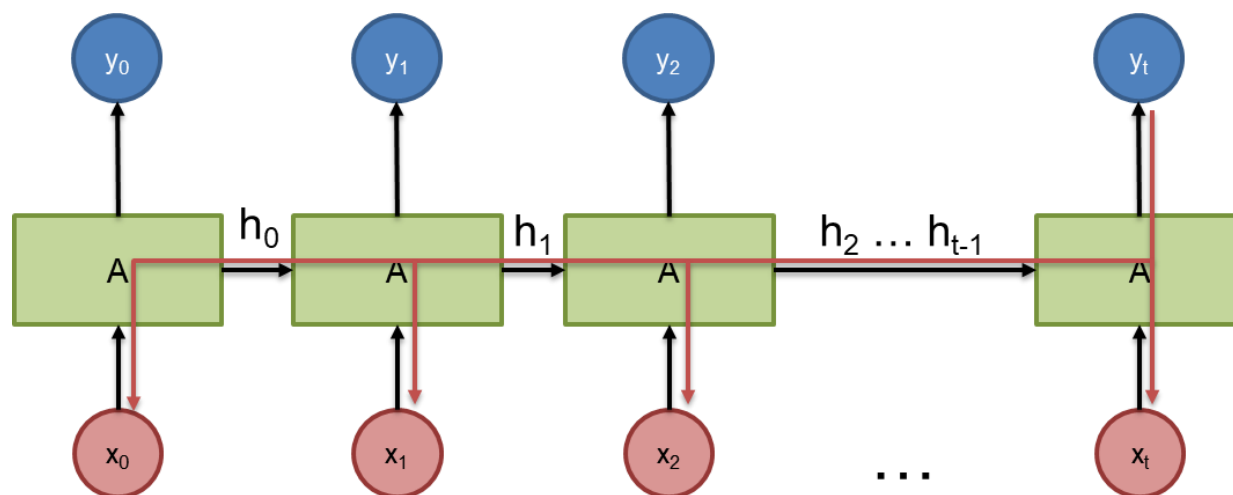
$$C(y, \hat{y}) = \sum_t C_t$$
$$C_t = \frac{1}{2} \|y_t - \hat{y}_t\|^2$$

Compute the gradient of the loss

$$\nabla C = [\nabla W_{xh}, \nabla W_{hh}, \nabla W_{hy}, \nabla b_h, \nabla b_y, \nabla h]$$

Backpropagation Through Time (BPTT)

MLTS Exercise 10



$$h_t = \tanh(W_{xh}x_t + W_{hh}h_{t-1} + b_h)$$

$$\tanh'(x) = 1 - \tanh(x)^2$$

$$\nabla W_{xh} = \tanh'(W_{hx}x_t + W_{hh}h_{t-1} + b_h) \cdot x_t^T$$

$$\nabla W_{hh} = \tanh'(W_{hx}x_t + W_{hh}h_{t-1} + b_h) \cdot h_{t-1}^T$$

$$\nabla b_h = \sum_{batch} \tanh'(W_{hx}x_t + W_{hh}h_{t-1} + b_h)$$

$$\nabla x_t = W_{hx}^T \cdot \tanh'(W_{hx}x_t + W_{hh}h_{t-1} + b_h)$$

$$\nabla h_{t-1} = W_{hh}^T \cdot \tanh'(W_{hx}x_{t-1} + W_{hh}h_{t-1} + b_h)$$

- Given the learning rate η update the weights and biases in the matrix W in the negative gradient direction


$$W = W - \eta \nabla C$$

- **Problem:** backpropagating through the whole sequence the parameter update is computational expensive
- **Solution:** truncated backpropagation through time (TBPTT)

Truncated Backpropagation Through Time

MLTS Exercise 10

- Instead of passing through the complete sequence, forward pass through a subset
- Backpropagate through the subset
- Continue with the next subset



Adapted from <https://tjmachinelearning.com/lectures/guest/rnnadv/rnnadv.html>



Long Short-Term Memory Network

Backward pass is completely linear:

- If weights are large (> 1), the gradient grows exponentially
- If weights are small (< 1), the gradients shrink exponentially

If we train on long sequences (> 100 time steps) the gradient may easily *explode* or *vanish*

The cat, which already ate ..., was/were full.

Backward pass is completely linear:


- If weights are large (> 1), the gradient grows exponentially
- If weights are small (< 1), the gradients shrink exponentially

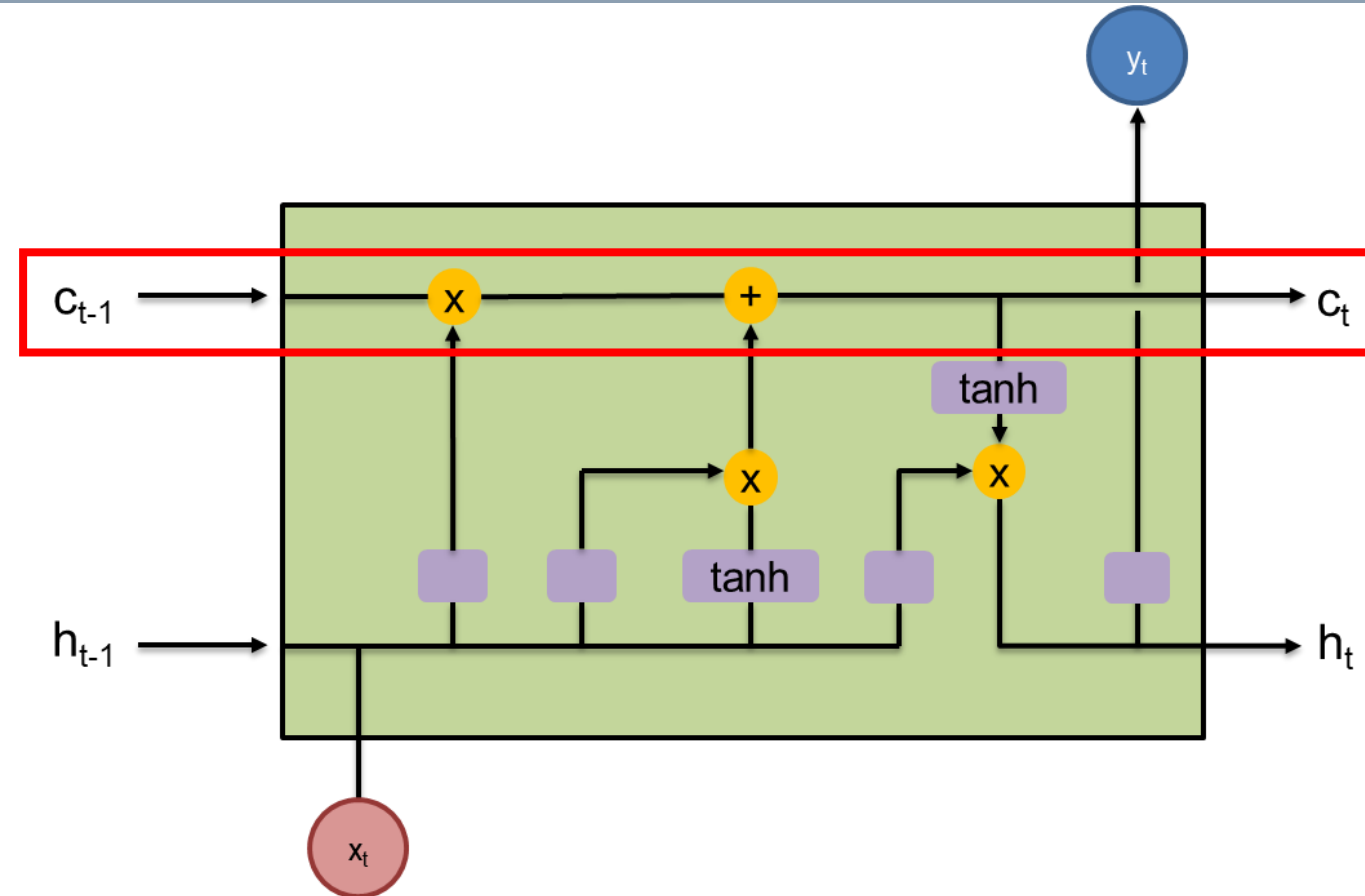
If we train on long sequences (> 100 time steps) the gradient may easily *explode* or *vanish*

Solution: Change RNN architecture

- e.g.: Long Short Term Memory Networks (LSTMs)

LSTM unit:



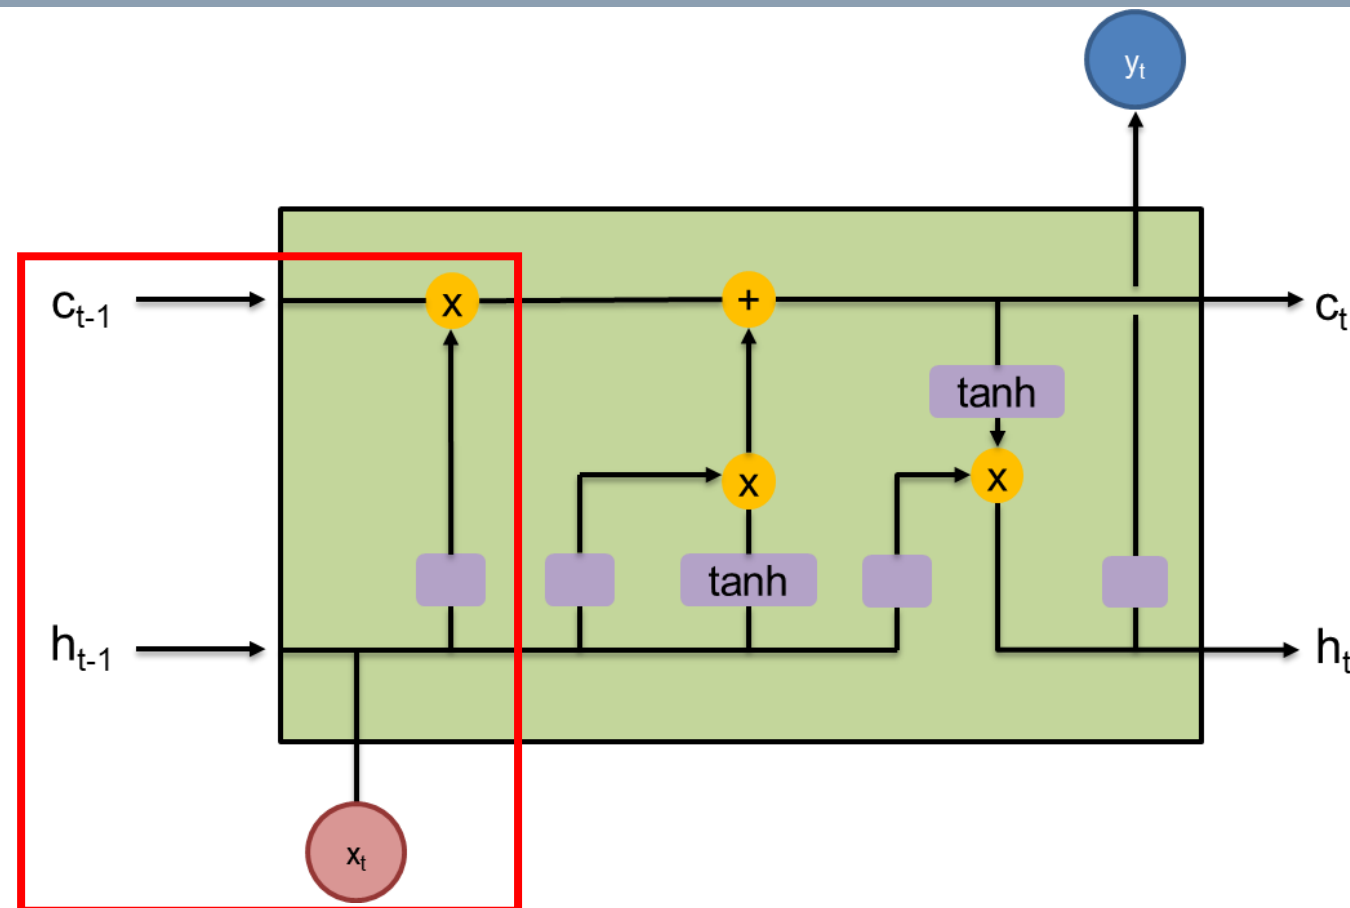


Cell state:

- Only minor linear interactions

- Easy flow of information, relatively unchanged

- Adding or removing information to the cell state is regulated by gates




Forget gate:

Decide how much of the previous cell state will be forgotten

Sigmoid layer squashes h_{t-1} and x_t between 0 and 1

$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$$

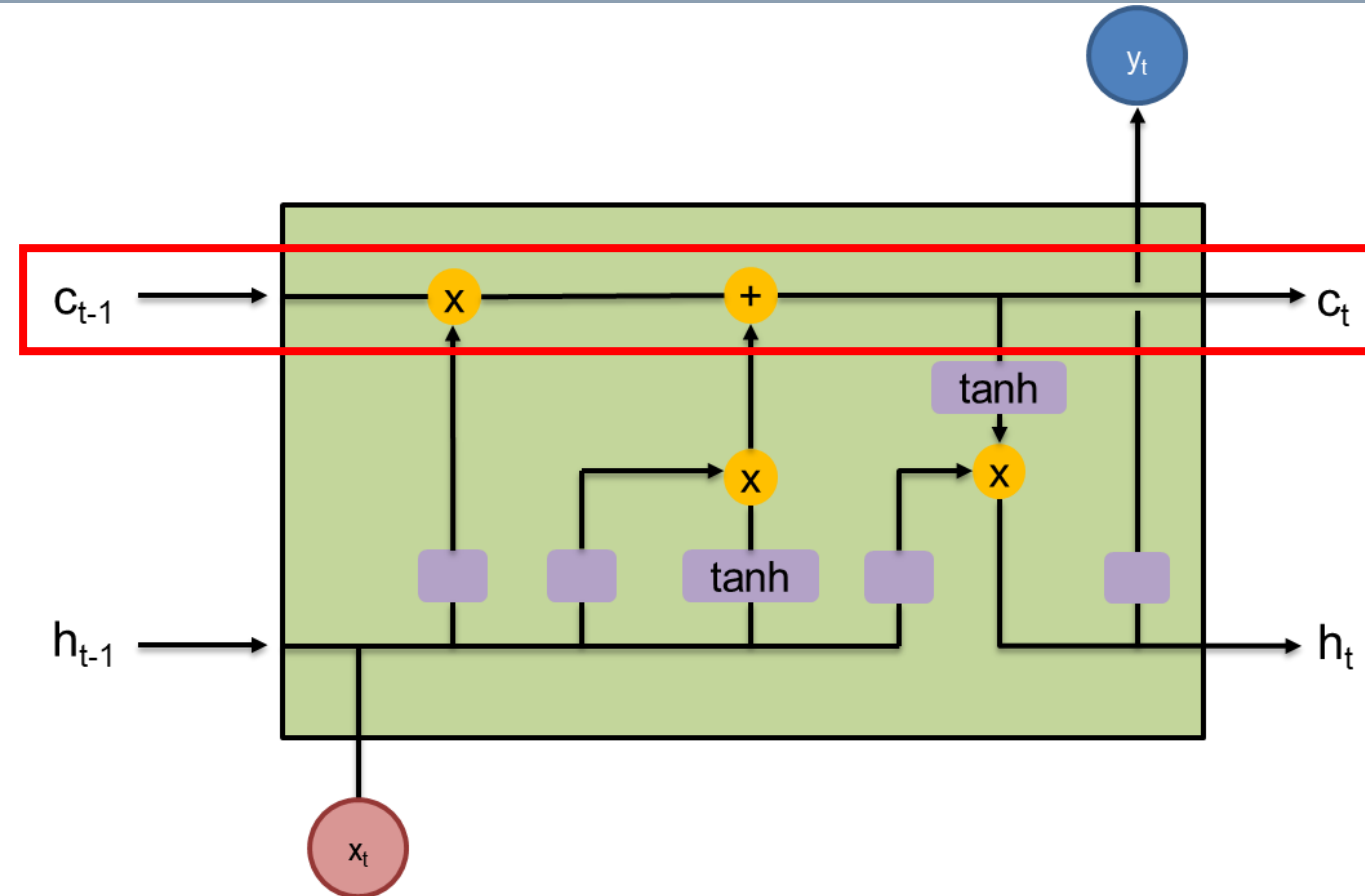


Input gate:

Decide what information we are going to store in the cell state

$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

$$g_t = \tanh(W_c \cdot [h_{t-1}, x_t] + b_c)$$




Combining values:

Update the old cell state c_{t-1} into c_t

$$c_t = f_t \odot c_{t-1} + i_t \odot g_t$$

Long Short-Term Memory Network

MLTS Exercise 10



Output gate:
Define output based on the cell state

$$o_t = \sigma(W_o \cdot [h_{t-1}, x_t] + b_o)$$

$$h_t = o_t \odot \tanh(c_t)$$

$$y_t = \sigma(W_y \cdot h_t + b_y)$$

LSTMs are great in order to avoid vanishing gradients

BUT: Loads of weights and biases which need to be optimized → difficult and slow training

Gated Recurrent Units (GRU) reduce the number of parameters to simplify training

Basically, hidden state and cell state are combined into a single parameter

- Forget and input gate are merged into an update gate

Concept was first presented by Cho et al. in 2014